

Copyright Notice

These slides are distributed under the Creative Commons License.

[DeepLearning.AI](#) makes these slides available for educational purposes. You may not use or distribute these slides for commercial purposes. You may make copies of these slides and use or distribute them for educational purposes as long as you cite [DeepLearning.AI](#) as the source of the slides.

For the rest of the details of the license, see <https://creativecommons.org/licenses/by-sa/2.0/legalcode>

audio	"imagenet2012"	text
"nsynth"	"imagenet2012_corrupted"	"cnn_dailymail"
	"kmnist"	"glue"
image	"lsun"	"imdb_reviews"
"abstract_reasoning"	"mnist"	"lm1b"
"caltech101"	"omniglot"	"multi_nli"
"cats_vs_dogs"	"open_images_v4"	"squad"
"celeb_a"	"oxford_iiit_pet"	"wikipedia"
"celeb_a_hq"	"quickdraw_bitmap"	"xnli"
"cifar10"	"rock_paper_scissors"	
"cifar100"	"shapes3d"	translate
"cifar10_corrupted"	"smallnorb"	"flores"
"coco2014"	"sun397"	"para_crawl"
"colorectal_histology"	"svhn_cropped"	"ted_hrlr_translate"
"cycle_gan"	"tf_flowers"	"ted_multi_translate"
"diabetic_retinopathy..."		"wmt15_translate"
"dsprites"	structured	"wmt16_translate"
"dtd"	"higgs"	"wmt17_translate"
"emnist"	"iris"	"wmt18_translate"
"fashion_mnist"	"titanic"	"wmt19_translate"
"horses_or_humans"		
"image_label_folder"		

audio	"imagenet2012"	text
"nsynth"	"imagenet2012_corrupted"	"cnn_dailymail"
	"kmnist"	"glue"
image	"lsun"	"imdb_reviews"
"abstract_reasoning"	"mnist"	"lm1b"
"caltech101"	"omniglot"	"multi_nli"
"cats_vs_dogs"	"open_images_v4"	"squad"
"celeb_a"	"oxford_iiit_pet"	"wikipedia"
"celeb_a_hq"	"quickdraw_bitmap"	"xnli"
"cifar10"	"rock_paper_scissors"	
"cifar100"	"shapes3d"	
"cifar10_corrupted"	"smallnorb"	translate
"coco2014"	"sun397"	"flores"
"colorectal_histology"	"svhn_cropped"	"para_crawl"
"cycle_gan"	"tf_flowers"	"ted_hrlr_translate"
"diabetic_retinopathy..."		"ted_multi_translate"
"dsprites"	structured	"wmt15_translate"
"dtd"	"higgs"	"wmt16_translate"
"emnist"	"iris"	"wmt17_translate"
"fashion_mnist"	"titanic"	"wmt18_translate"
"horses_or_humans"		"wmt19_translate"
"image_label_folder"		

<http://ai.stanford.edu/~amaas/data/sentiment/>

```
@InProceedings{maas-EtAl:2011:ACL-HLT2011,  
  author      = {Maas, Andrew L. and Daly, Raymond E. and Pham, Peter T. and Huang, Dan and Ng,  
Andrew Y. and Potts, Christopher},  
  title       = {Learning Word Vectors for Sentiment Analysis},  
  booktitle   = {Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics:  
Human Language Technologies},  
  month       = {June},  
  year        = {2011},  
  address     = {Portland, Oregon, USA},  
  publisher   = {Association for Computational Linguistics},  
  pages       = {142--150},  
  url         = {http://www.aclweb.org/anthology/P11-1015}  
}
```

```
import tensorflow_datasets as tfds
```

```
imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)
```

```
import tensorflow_datasets as tfds
```

```
imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)
```

```
single_example = list(imdb['train'].take(1))[0]
```

```
print(single_example[0])
```

```
tf.Tensor(b"This was an absolutely terrible movie. Don't be lured in by Christopher Walken  
or Michael Ironside. Both are great actors, but this must simply be their worst role in  
history. Even their great acting could not redeem this movie's ridiculous storyline. This  
movie is an early nineties US propaganda piece. The most pathetic scenes were those when  
the Columbian rebels were making their cases for revolutions. Maria Conchita Alonso  
appeared phony, and her pseudo-love affair with Walken was nothing but a pathetic  
emotional plug in a movie that was devoid of any real meaning. I am disappointed that  
there are movies like this, ruining actor's like Christopher Walken's good name. I could  
barely sit through it.", shape=(), dtype=string)
```

```
import tensorflow_datasets as tfds
```

```
imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)
```

```
single_example = list(imdb['train'].take(1))[0]
```

```
print(single_example[0])
```

```
tf.Tensor(b"This was an absolutely terrible movie. Don't be lured in by Christopher Walken  
or Michael Ironside. Both are great actors, but this must simply be their worst role in  
history. Even their great acting could not redeem this movie's ridiculous storyline. This  
movie is an early nineties US propaganda piece. The most pathetic scenes were those when  
the Columbian rebels were making their cases for revolutions. Maria Conchita Alonso  
appeared phony, and her pseudo-love affair with Walken was nothing but a pathetic  
emotional plug in a movie that was devoid of any real meaning. I am disappointed that  
there are movies like this, ruining actor's like Christopher Walken's good name. I could  
barely sit through it.", shape=(), dtype=string)
```

```
import tensorflow_datasets as tfds

imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)

single_example = list(imdb['train'].take(1))[0]

print(single_example[1])

tf.Tensor(0, shape=(), dtype=int64)
```



```
import tensorflow_datasets as tfds
```

```
imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)
```

```
train_data, test_data = imdb['train'], imdb['test']
```

```
train_reviews = train_dataset.map(lambda review, label: review)
```

```
train_labels = train_dataset.map(lambda review, label: label)
```

```
test_reviews = test_dataset.map(lambda review, label: review)
```

```
test_labels = test_dataset.map(lambda review, label: label)
```

```
import tensorflow_datasets as tfds

imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)

train_data, test_data = imdb['train'], imdb['test']

train_reviews = train_dataset.map(lambda review, label: review)
train_labels = train_dataset.map(lambda review, label: label)

test_reviews = test_dataset.map(lambda review, label: review)
test_labels = test_dataset.map(lambda review, label: label)
```

```
import tensorflow_datasets as tfds

imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)

train_data, test_data = imdb['train'], imdb['test']

train_reviews = train_dataset.map(lambda review, label: review)
train_labels = train_dataset.map(lambda review, label: label)

test_reviews = test_dataset.map(lambda review, label: review)
test_labels = test_dataset.map(lambda review, label: label)
```

```
vectorize_layer = tf.keras.layers.TextVectorization(max_tokens=10000)

vectorize_layer.adapt(train_reviews)

def padding_func(sequences):
    sequences = sequences.ragged_batch(batch_size=sequences.cardinality())
    sequences = sequences.get_single_element()

    padded_sequences = tf.keras.utils.pad_sequences(sequences.numpy(), maxlen=120,
                                                    truncating='post', padding='pre')

    padded_sequences = tf.data.Dataset.from_tensor_slices(padded_sequences)

    return padded_sequences

train_sequences = train_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
test_sequences = test_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
```



```
vectorize_layer = tf.keras.layers.TextVectorization(max_tokens=10000)
```

```
vectorize_layer.adapt(train_reviews)
```

```
def padding_func(sequences):  
    sequences = sequences.ragged_batch(batch_size=sequences.cardinality())  
    sequences = sequences.get_single_element()  
  
    padded_sequences = tf.keras.utils.pad_sequences(sequences.numpy(), maxlen=120,  
                                                    truncating='post', padding='pre')  
  
    padded_sequences = tf.data.Dataset.from_tensor_slices(padded_sequences)  
  
    return padded_sequences
```

```
train_sequences = train_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)  
test_sequences = test_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
```



```
vectorize_layer = tf.keras.layers.TextVectorization(max_tokens=10000)

vectorize_layer.adapt(train_reviews)

def padding_func(sequences):
    sequences = sequences.ragged_batch(batch_size=sequences.cardinality())
    sequences = sequences.get_single_element()

    padded_sequences = tf.keras.utils.pad_sequences(sequences.numpy(), maxlen=120,
                                                    truncating='post', padding='pre')

    padded_sequences = tf.data.Dataset.from_tensor_slices(padded_sequences)

    return padded_sequences

train_sequences = train_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
test_sequences = test_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
```



```
vectorize_layer = tf.keras.layers.TextVectorization(max_tokens=10000)
```

```
vectorize_layer.adapt(train_reviews)
```

```
def padding_func(sequences):  
    sequences = sequences.ragged_batch(batch_size=sequences.cardinality())  
    sequences = sequences.get_single_element()  
  
    padded_sequences = tf.keras.utils.pad_sequences(sequences.numpy(), maxlen=120,  
                                                    truncating='post', padding='pre')  
  
    padded_sequences = tf.data.Dataset.from_tensor_slices(padded_sequences)  
  
    return padded_sequences
```

```
train_sequences = train_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)  
test_sequences = test_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
```



```
vectorize_layer = tf.keras.layers.TextVectorization(max_tokens=10000)
```

```
vectorize_layer.adapt(train_reviews)
```

```
def padding_func(sequences):  
    sequences = sequences.ragged_batch(batch_size=sequences.cardinality())  
    sequences = sequences.get_single_element()  
  
    padded_sequences = tf.keras.utils.pad_sequences(sequences.numpy(), maxlen=120,  
                                                    truncating='post', padding='pre')  
  
    padded_sequences = tf.data.Dataset.from_tensor_slices(padded_sequences)  
  
    return padded_sequences
```

```
train_sequences = train_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
```

```
test_sequences = test_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
```



```
vectorize_layer = tf.keras.layers.TextVectorization(max_tokens=10000)

vectorize_layer.adapt(train_reviews)

def padding_func(sequences):
    sequences = sequences.ragged_batch(batch_size=sequences.cardinality())
    sequences = sequences.get_single_element()

    padded_sequences = tf.keras.utils.pad_sequences(sequences.numpy(), maxlen=120,
                                                    truncating='post', padding='pre')

    padded_sequences = tf.data.Dataset.from_tensor_slices(padded_sequences)

    return padded_sequences

train_sequences = train_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
test_sequences = test_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
```



```
vectorize_layer = tf.keras.layers.TextVectorization(max_tokens=10000)

vectorize_layer.adapt(train_reviews)

def padding_func(sequences):
    sequences = sequences.ragged_batch(batch_size=sequences.cardinality())
    sequences = sequences.get_single_element()

    padded_sequences = tf.keras.utils.pad_sequences(sequences.numpy(), maxlen=120,
                                                    truncating='post', padding='pre')

    padded_sequences = tf.data.Dataset.from_tensor_slices(padded_sequences)

    return padded_sequences

train_sequences = train_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
test_sequences = test_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
```

```
vectorize_layer = tf.keras.layers.TextVectorization(max_tokens=10000)

vectorize_layer.adapt(train_reviews)

def padding_func(sequences):
    sequences = sequences.ragged_batch(batch_size=sequences.cardinality())
    sequences = sequences.get_single_element()

    padded_sequences = tf.keras.utils.pad_sequences(sequences.numpy(), maxlen=120,
                                                    truncating='post', padding='pre')

    padded_sequences = tf.data.Dataset.from_tensor_slices(padded_sequences)

    return padded_sequences

train_sequences = train_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
test_sequences = test_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
```



```
vectorize_layer = tf.keras.layers.TextVectorization(max_tokens=10000)

vectorize_layer.adapt(train_reviews)

def padding_func(sequences):
    sequences = sequences.ragged_batch(batch_size=sequences.cardinality())
    sequences = sequences.get_single_element()

    padded_sequences = tf.keras.utils.pad_sequences(sequences.numpy(), maxlen=120,
                                                    truncating='post', padding='pre')

    padded_sequences = tf.data.Dataset.from_tensor_slices(padded_sequences)

    return padded_sequences

train_sequences = train_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
test_sequences = test_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
```



```
vectorize_layer = tf.keras.layers.TextVectorization(max_tokens=10000)

vectorize_layer.adapt(train_reviews)

def padding_func(sequences):
    sequences = sequences.ragged_batch(batch_size=sequences.cardinality())
    sequences = sequences.get_single_element()

    padded_sequences = tf.keras.utils.pad_sequences(sequences.numpy(), maxlen=120,
                                                    truncating='post', padding='pre')

    padded_sequences = tf.data.Dataset.from_tensor_slices(padded_sequences)

    return padded_sequences

train_sequences = train_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
test_sequences = test_reviews.map(lambda text: vectorize_layer(text)).apply(padding_func)
```



```
train_dataset_vectorized = tf.data.Dataset.zip(train_sequences, train_labels)
test_dataset_vectorized = tf.data.Dataset.zip(test_sequences, test_labels)

SHUFFLE_BUFFER_SIZE = 1000
PREFETCH_BUFFER_SIZE = tf.data.AUTOTUNE
BATCH_SIZE = 32

train_dataset_final = (train_dataset_vectorized
    .cache()
    .shuffle(SHUFFLE_BUFFER_SIZE)
    .prefetch(PREFETCH_BUFFER_SIZE)
    .batch(BATCH_SIZE)
)

test_dataset_final = (test_dataset_vectorized
    .cache()
    .prefetch(PREFETCH_BUFFER_SIZE)
    .batch(BATCH_SIZE)
)
```

```
train_dataset_vectorized = tf.data.Dataset.zip(train_sequences, train_labels)
test_dataset_vectorized = tf.data.Dataset.zip(test_sequences, test_labels)
```

```
SHUFFLE_BUFFER_SIZE = 1000
PREFETCH_BUFFER_SIZE = tf.data.AUTOTUNE
BATCH_SIZE = 32
```

```
train_dataset_final = (train_dataset_vectorized
                        .cache()
                        .shuffle(SHUFFLE_BUFFER_SIZE)
                        .prefetch(PREFETCH_BUFFER_SIZE)
                        .batch(BATCH_SIZE)
                        )
```

```
test_dataset_final = (test_dataset_vectorized
                       .cache()
                       .prefetch(PREFETCH_BUFFER_SIZE)
                       .batch(BATCH_SIZE)
                       )
```

```
train_dataset_vectorized = tf.data.Dataset.zip(train_sequences, train_labels)
test_dataset_vectorized = tf.data.Dataset.zip(test_sequences, test_labels)
```

```
SHUFFLE_BUFFER_SIZE = 1000
PREFETCH_BUFFER_SIZE = tf.data.AUTOTUNE
BATCH_SIZE = 32

train_dataset_final = (train_dataset_vectorized
                        .cache()
                        .shuffle(SHUFFLE_BUFFER_SIZE)
                        .prefetch(PREFETCH_BUFFER_SIZE)
                        .batch(BATCH_SIZE)
                        )

test_dataset_final = (test_dataset_vectorized
                      .cache()
                      .prefetch(PREFETCH_BUFFER_SIZE)
                      .batch(BATCH_SIZE)
                      )
```



```
model = tf.keras.Sequential([  
    tf.keras.Input(shape=(120,)),  
    tf.keras.layers.Embedding(vocab_size, embedding_dim),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(6, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```

```
model = tf.keras.Sequential([  
    tf.keras.Input(shape=(120,)),  
    tf.keras.layers.Embedding(vocab_size, embedding_dim),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(6, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```

```
model = tf.keras.Sequential([  
    tf.keras.Input(shape=(120,)),  
    tf.keras.layers.Embedding(vocab_size, embedding_dim),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(6, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```

Layer (type)	Output Shape	Param #
embedding_9 (Embedding)	(None, 120, 16)	160000
flatten_3 (Flatten)	(None, 1920)	0
dense_14 (Dense)	(None, 6)	11526
dense_15 (Dense)	(None, 1)	7
Total params: 171,533		
Trainable params: 171,533		
Non-trainable params: 0		

```
model = tf.keras.Sequential([  
    tf.keras.Input(shape=(120,)),  
    tf.keras.layers.Embedding(vocab_size, embedding_dim),  
    tf.keras.layers.GlobalAveragePooling1D(),  
    tf.keras.layers.Dense(6, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```



Layer (type)	Output Shape	Param #
embedding_11 (Embedding)	(None, 120, 16)	160000
global_average_pooling1d_3 ((None, 16)	0
dense_16 (Dense)	(None, 6)	102
dense_17 (Dense)	(None, 1)	7
Total params: 160,109		
Trainable params: 160,109		
Non-trainable params: 0		

```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])  
model.summary()
```

```
num_epochs = 10  
model.fit(train_dataset_final,  
          epochs=num_epochs,  
          validation_data=test_dataset_final)
```


Epoch 8/10

25000/25000 [=====] -

6s 256us/sample - loss: 5.2086e-04 - acc: 1.0000 - val_loss: 0.7252 - val_acc: 0.8270

Epoch 9/10

25000/25000 [=====] -

6s 222us/sample - loss: 3.0199e-04 - acc: 1.0000 - val_loss: 0.7628 - val_acc: 0.8269

Epoch 10/10

25000/25000 [=====] -

6s 224us/sample - loss: 1.7872e-04 - acc: 1.0000 - val_loss: 0.7997 - val_acc: 0.8259

```
embedding_layer = model.layers[0]  
embedding_weights = embedding_layer.get_weights()[0]  
print(embedding_weights.shape) # shape: (vocab_size, embedding_dim)
```

```
(10000, 16)
```

```
import io

out_v = io.open('vecs.tsv', 'w', encoding='utf-8')
out_m = io.open('meta.tsv', 'w', encoding='utf-8')

vocabulary = vectorize_layer.get_vocabulary()

for word_num in range(1, len(vocabulary)):
    word_name = vocabulary[word_num]
    word_embedding = embedding_weights[word_num]
    out_m.write(word_name + "\n")
    out_v.write('\t'.join([str(x) for x in word_embedding]) + "\n")

out_v.close()
out_m.close()
```

```
import io

out_v = io.open('vecs.tsv', 'w', encoding='utf-8')
out_m = io.open('meta.tsv', 'w', encoding='utf-8')

vocabulary = vectorize_layer.get_vocabulary()

for word_num in range(1, len(vocabulary)):
    word_name = vocabulary[word_num]
    word_embedding = embedding_weights[word_num]
    out_m.write(word_name + "\n")
    out_v.write('\t'.join([str(x) for x in word_embedding]) + "\n")

out_v.close()
out_m.close()
```

```
import io
```

```
out_v = io.open('vecs.tsv', 'w', encoding='utf-8')
```

```
out_m = io.open('meta.tsv', 'w', encoding='utf-8')
```

```
vocabulary = vectorize_layer.get_vocabulary()
```

```
for word_num in range(1, len(vocabulary)):
    word_name = vocabulary[word_num]
    word_embedding = embedding_weights[word_num]
    out_m.write(word_name + "\n")
    out_v.write('\t'.join([str(x) for x in word_embedding]) + "\n")
```

```
out_v.close()
```

```
out_m.close()
```

```
import io

out_v = io.open('vecs.tsv', 'w', encoding='utf-8')
out_m = io.open('meta.tsv', 'w', encoding='utf-8')

vocabulary = vectorize_layer.get_vocabulary()

for word_num in range(1, len(vocabulary)):
    word_name = vocabulary[word_num]
    word_embedding = embedding_weights[word_num]
    out_m.write(word_name + '\n')
    out_v.write('\t'.join([str(x) for x in word_embedding]) + '\n')

out_v.close()
out_m.close()
```

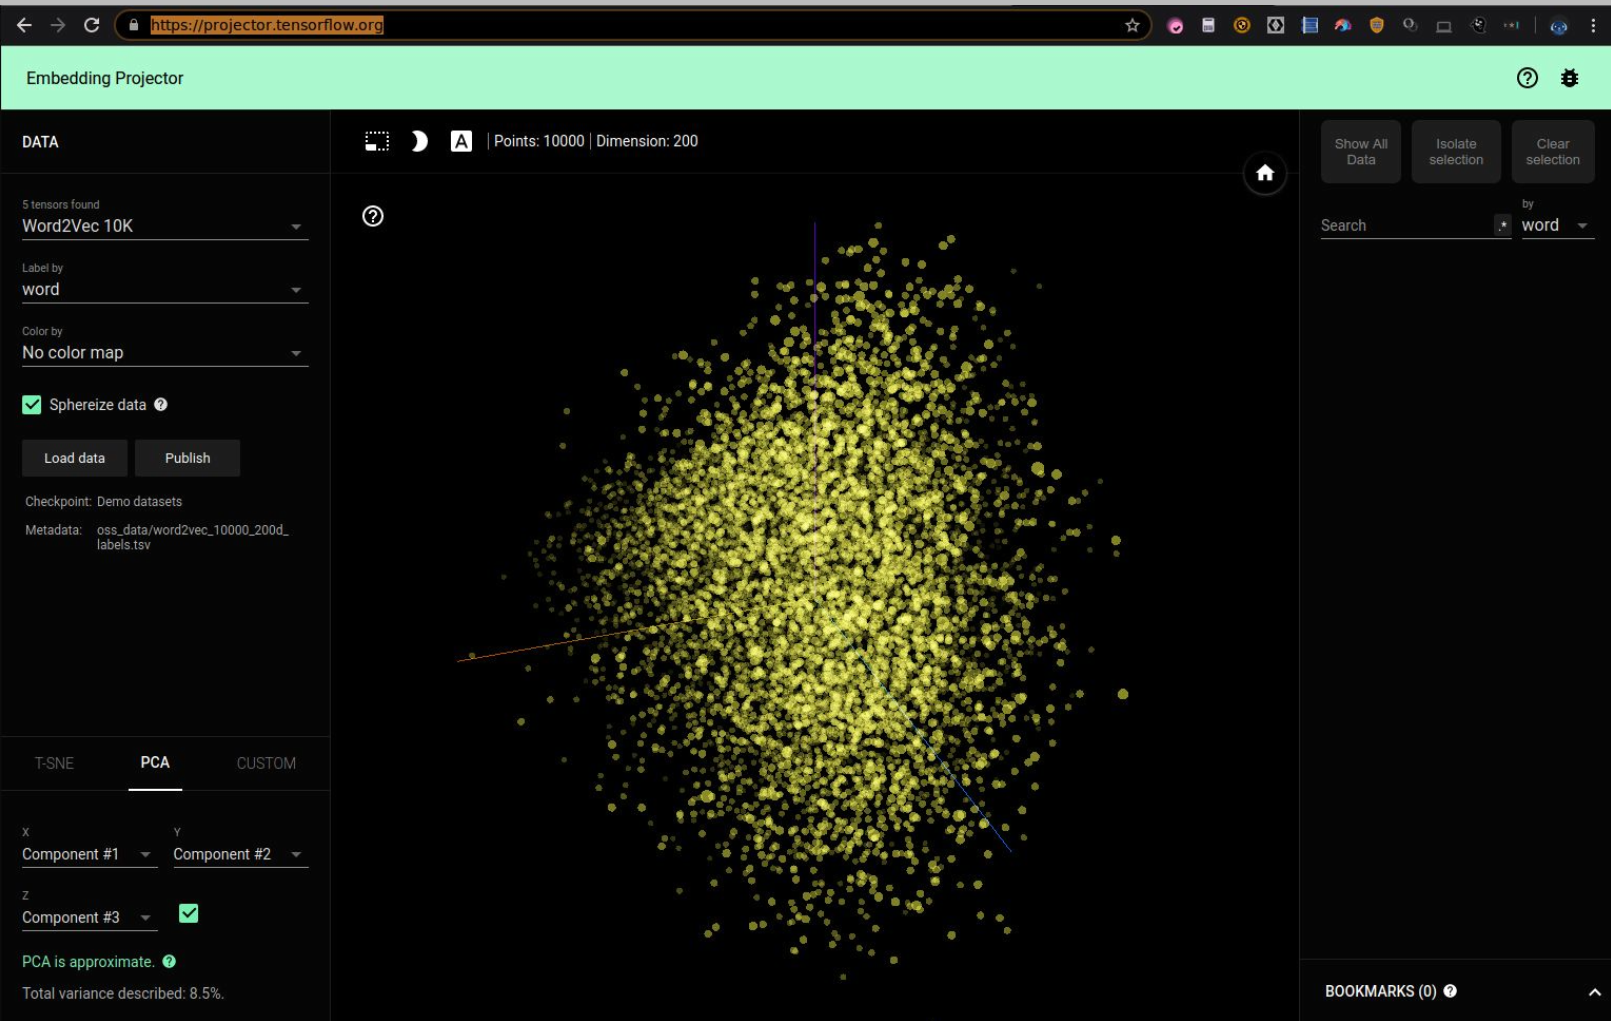
```
import io

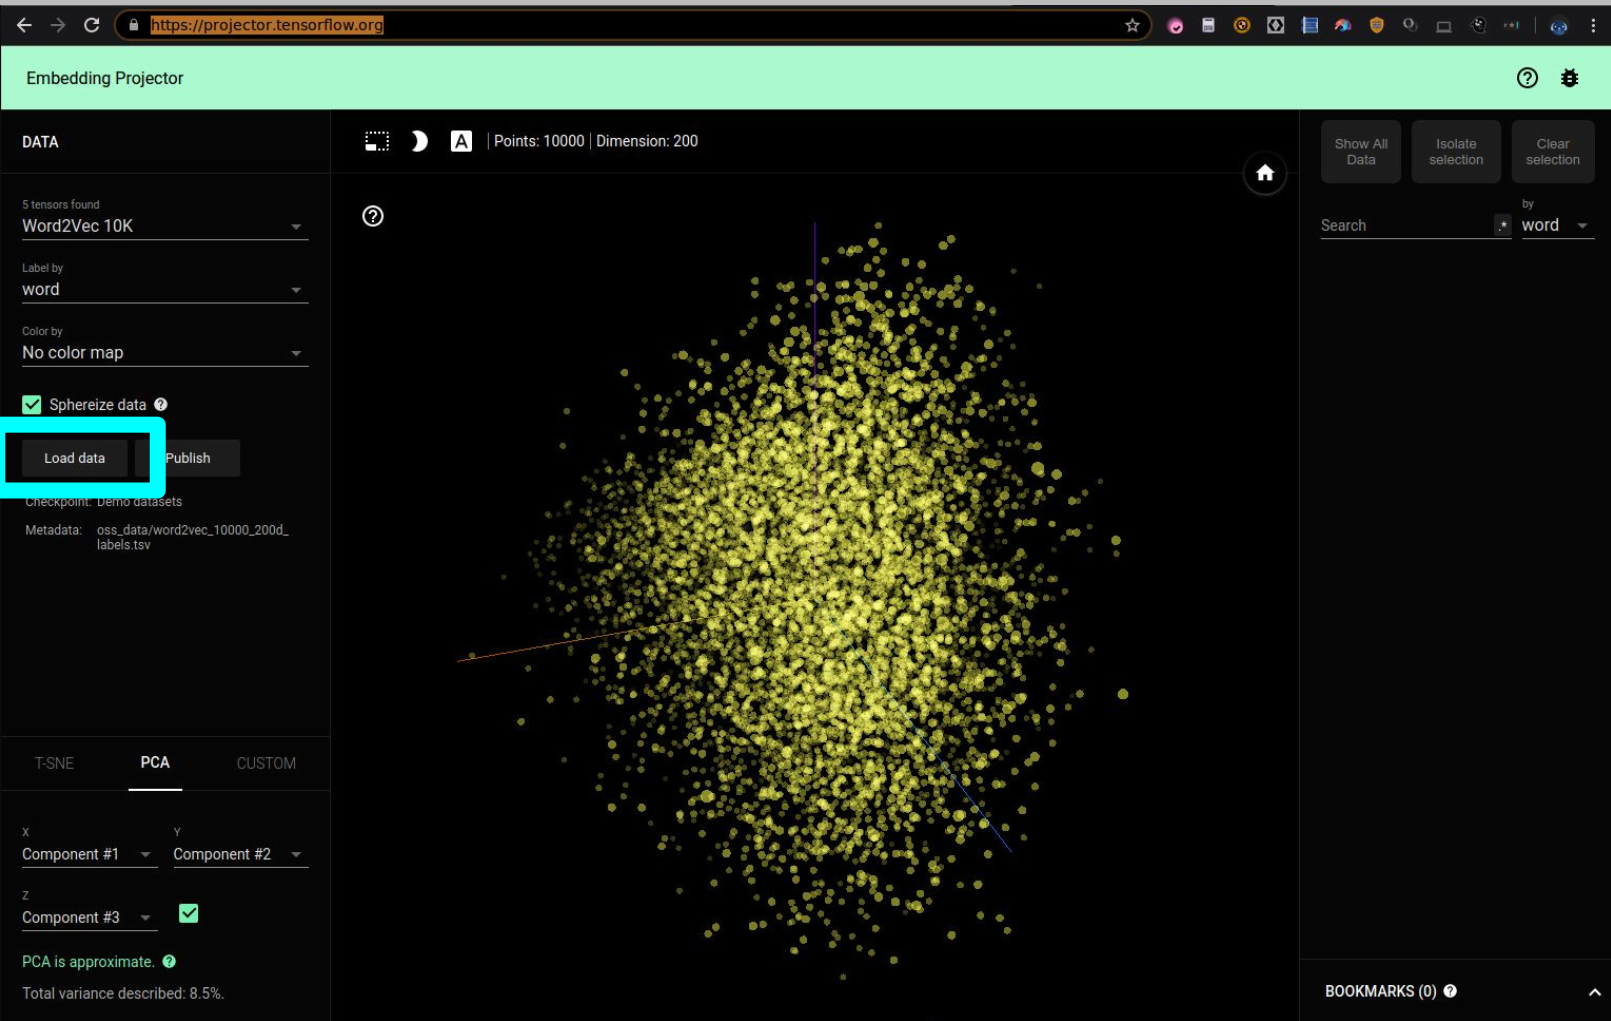
out_v = io.open('vecs.tsv', 'w', encoding='utf-8')
out_m = io.open('meta.tsv', 'w', encoding='utf-8')

vocabulary = vectorize_layer.get_vocabulary()

for word_num in range(1, len(vocabulary)):
    word_name = vocabulary[word_num]
    word_embedding = embedding_weights[word_num]
    out_m.write(word_name + "\n")
    out_v.write('\t'.join([str(x) for x in word_embedding]) + "\n")

out_v.close()
out_m.close()
```





Load data from your computer

Step 1: Load a TSV file of vectors.

Example of 3 vectors with dimension 4:

```
0.1\t0.2\t0.5\t0.9  
0.2\t0.1\t5.0\t0.2  
0.4\t0.1\t7.0\t0.8
```

Choose file

Step 2 (optional): Load a TSV file of metadata.

Example of 3 data points and 2 columns.

Note: If there is more than one column, the first row will be parsed as column labels.

```
Pokémon\tSpecies  
Wartortle\tTurtle  
Venusaur\tSeed  
Charmeleon\tFlame
```

Choose file

Click outside to dismiss.

DATA

5 tensors found

Word2Vec 10K

☐ Sphereize data ?

Load data

Publish

Checkpoint: vecs.tsv

Metadata: meta.tsv

T-SNE

PCA

CUSTOM

X
Component #1

Y
Component #2

Z
Component #3



PCA is approximate. ?

Total variance described: 98.7%.



Points: 9999 | Dimension: 16



Show All
Data

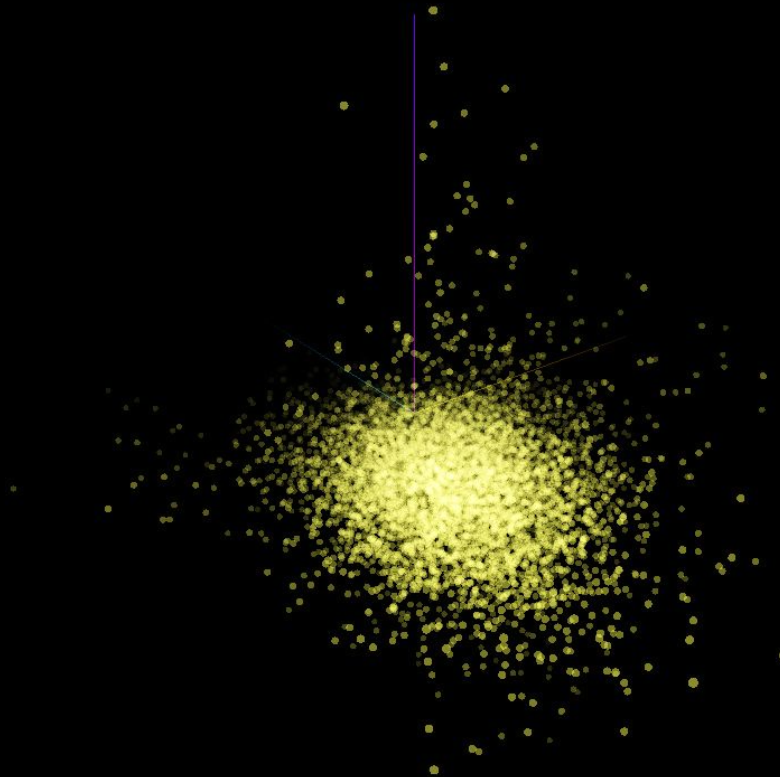
Isolate
selection

Clear
selection

Search

.*

by




BOOKMARKS (0) ?



DATA

5 tensors found

Word2Vec 10K

☒ Sphereize data 

Load data

Publish

Checkpoint: vecs.tsv

Metadata: meta.tsv

T-SNE

PCA

CUSTOM

x


Component #1

y

Component #2

z

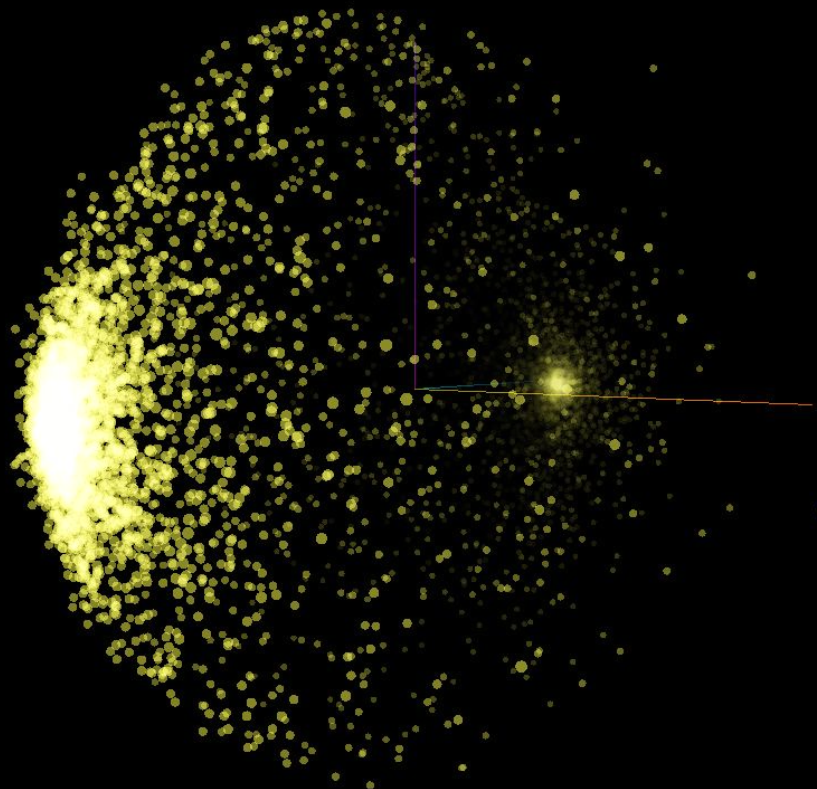
Component #3

PCA is approximate. 

Total variance described: 88.7%.




Points: 9999 | Dimension: 16

Show All
DataIsolate
selectionClear
selection

Search



by

BOOKMARKS (0) 

TRAINING_SIZE = 20000

VOCAB_SIZE = 10000

MAX_LENGTH = 32

EMBEDDING_DIM = 16

```
with open("/tmp/sarcasm.json", 'r') as f:  
    datastore = json.load(f)  
  
sentences = []  
labels = []  
  
for item in datastore:  
    sentences.append(item['headline'])  
    labels.append(item['is_sarcastic'])
```

```
training_sentences = sentences[0:training_size]
testing_sentences = sentences[training_size:]
training_labels = labels[0:training_size]
testing_labels = labels[training_size:]
```

```
training_sentences = sentences[0:training_size]  
testing_sentences = sentences[training_size:]  
training_labels = labels[0:training_size]  
testing_labels = labels[training_size:]
```



```
training_sentences = sentences[0:training_size]  
testing_sentences = sentences[training_size:]  
training_labels = labels[0:training_size]  
testing_labels = labels[training_size:]
```

```
training_sentences = sentences[0:training_size]
testing_sentences = sentences[training_size:]
training_labels = labels[0:training_size]
testing_labels = labels[training_size:]
```

```
vectorize_layer = tf.keras.layers.TextVectorization(  
    max_tokens=VOCAB_SIZE,  
    output_sequence_length=MAX_LENGTH)  
  
vectorize_layer.adapt(train_sentences)  
  
train_sequences = vectorize_layer(train_sentences)  
test_sequences = vectorize_layer(test_sentences)  
  
train_dataset_vectorized = tf.data.Dataset.from_tensor_slices(  
    (train_sequences, train_labels))  
test_dataset_vectorized = tf.data.Dataset.from_tensor_slices(  
    (test_sequences, test_labels))
```

```
vectorize_layer = tf.keras.layers.TextVectorization(  
    max_tokens=VOCAB_SIZE,  
    output_sequence_length=MAX_LENGTH)  
  
vectorize_layer.adapt(train_sentences)
```

```
train_sequences = vectorize_layer(train_sentences)  
test_sequences = vectorize_layer(test_sentences)
```

```
train_dataset_vectorized = tf.data.Dataset.from_tensor_slices(  
    (train_sequences, train_labels))  
test_dataset_vectorized = tf.data.Dataset.from_tensor_slices(  
    (test_sequences, test_labels))
```

```
vectorize_layer = tf.keras.layers.TextVectorization(  
    max_tokens=VOCAB_SIZE,  
    output_sequence_length=MAX_LENGTH)  
  
vectorize_layer.adapt(train_sentences)  
  
train_sequences = vectorize_layer(train_sentences)  
test_sequences = vectorize_layer(test_sentences)  
  
train_dataset_vectorized = tf.data.Dataset.from_tensor_slices(  
    (train_sequences, train_labels))  
test_dataset_vectorized = tf.data.Dataset.from_tensor_slices(  
    (test_sequences, test_labels))
```

```
vectorize_layer = tf.keras.layers.TextVectorization(  
    max_tokens=VOCAB_SIZE,  
    output_sequence_length=MAX_LENGTH)
```

```
vectorize_layer.adapt(train_sentences)
```

```
train_sequences = vectorize_layer(train_sentences)
```

```
test_sequences = vectorize_layer(test_sentences)
```

```
train_dataset_vectorized = tf.data.Dataset.from_tensor_slices(  
    (train_sequences, train_labels))  
test_dataset_vectorized = tf.data.Dataset.from_tensor_slices(  
    (test_sequences, test_labels))
```



```
SHUFFLE_BUFFER_SIZE = 1000
PREFETCH_BUFFER_SIZE = tf.data.AUTOTUNE
BATCH_SIZE = 32

train_dataset_final = (train_dataset_vectorized
    .cache()
    .shuffle(SHUFFLE_BUFFER_SIZE)
    .prefetch(PREFETCH_BUFFER_SIZE)
    .batch(BATCH_SIZE)
)

test_dataset_final = (test_dataset_vectorized
    .cache()
    .prefetch(PREFETCH_BUFFER_SIZE)
    .batch(BATCH_SIZE)
)
```

```
model = tf.keras.Sequential([
    tf.keras.Input(shape=(MAX_LENGTH,)),
    tf.keras.layers.Embedding(VOCAB_SIZE, EMBEDDING_DIM),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(6, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```




```
model.summary()
```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 32, 16)	160000
global_average_pooling1d_2 ((None, 16)		0
dense_4 (Dense)	(None, 24)	408
dense_5 (Dense)	(None, 1)	25
Total params: 160,433		
Trainable params: 160,433		
Non-trainable params: 0		

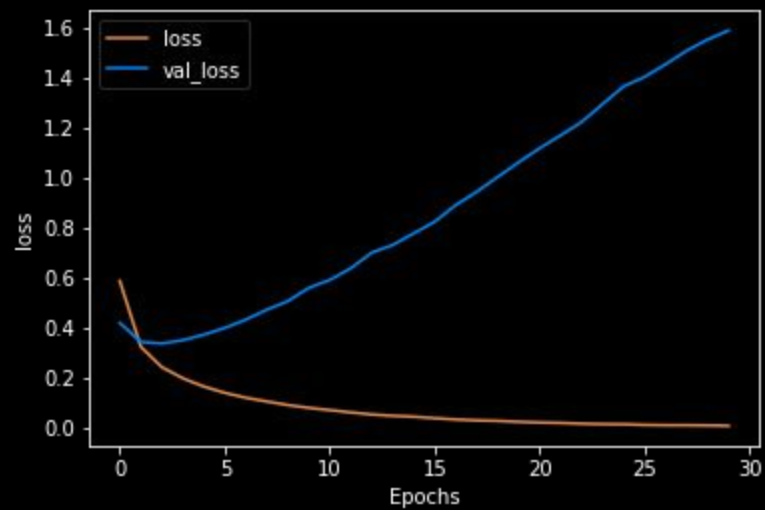
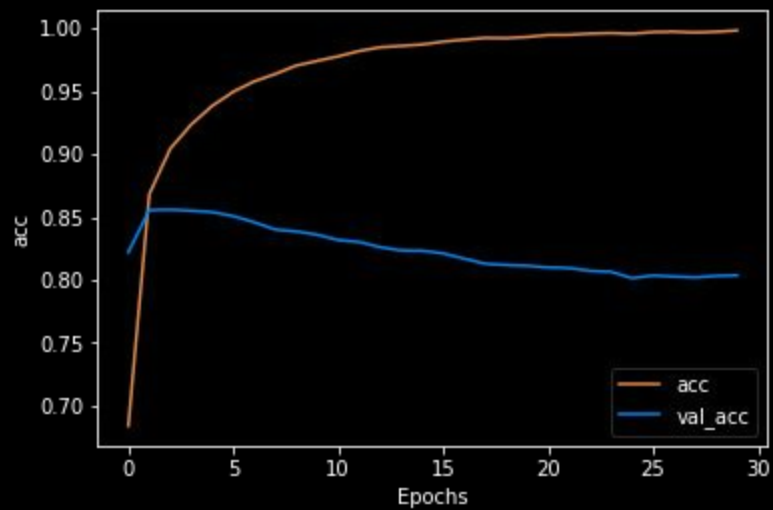
```
num_epochs = 30
```

```
history = model.fit(train_dataset_final, epochs=num_epochs,  
                    validation_data=test_dataset_final, verbose=2)
```

```
import matplotlib.pyplot as plt

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_' + string])
    plt.xlabel("Epochs")
    plt.ylabel(string)
    plt.legend([string, 'val_' + string])
    plt.show()

plot_graphs(history, "accuracy")
plot_graphs(history, "loss")
```


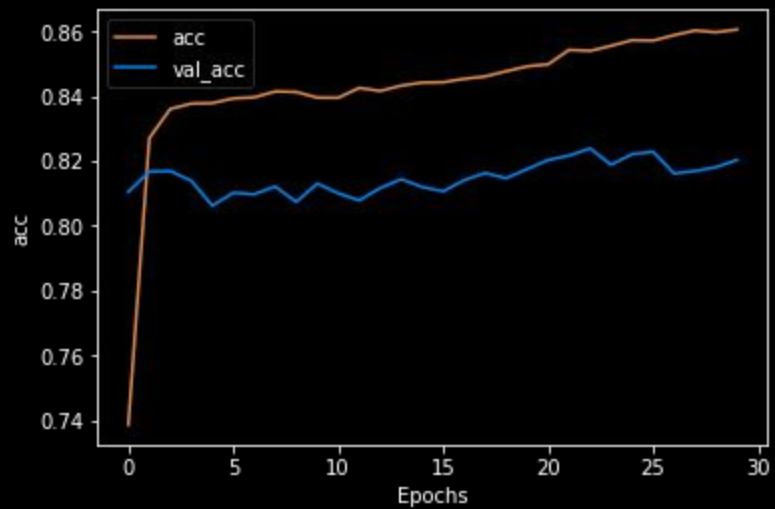


TRAINING_SIZE = 20000

VOCAB_SIZE = 1000 (was 10000)

MAX_LENGTH = 16 (was 32)

EMBEDDING_DIM = 16


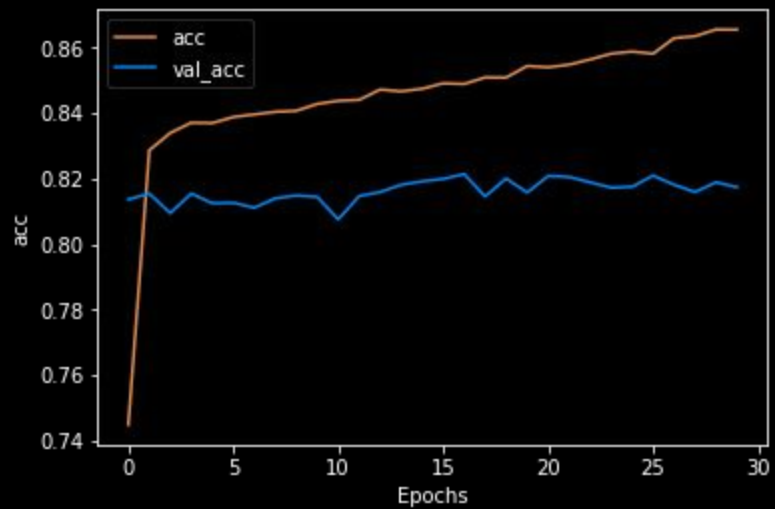


TRAINING_SIZE = 20000

VOCAB_SIZE = 1000 (was 10000)


MAX_LENGTH = 16 (was 32)

EMBEDDING_DIM = 32 (was 16)




Word tokenization

I love NLP




Word tokenization

I love NLP



Subword tokenization

I love NLP



<https://github.com/tensorflow/datasets/tree/master/docs/catalog>

275 lines (209 sloc) | 9.65 KB

<> [icon] Raw Blame [icon] [icon] [icon]

imdb_reviews

- Description:

Large Movie Review Dataset. This is a dataset for binary sentiment classification containing substantially more data than previous benchmark datasets. We provide a set of 25,000 highly polar movie reviews for training, and 25,000 for testing. There is additional unlabeled data for use as well.

- Homepage: <http://ai.stanford.edu/~amaas/data/sentiment/>

- Source code: `tfds.text.IMDBReviews`

- Versions:

- `1.0.0` (default): New split API (<https://tensorflow.org/datasets/splits>)

- Download size: `80.23 MiB`

- Dataset size: `Unknown size`

- Auto-cached ([documentation](#)): Unknown

- Splits:

Split	Examples
<code>'test'</code>	25,000
<code>'train'</code>	25,000
<code>'unsupervised'</code>	50,000

275 lines (209 sloc) | 9.65 KB

<> [icon] Raw Blame [icon] [icon] [icon]

imdb_reviews/plain_text (default config)

- Config description: Plain text
- Features:

```
FeaturesDict({  
  'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=2),  
  'text': Text(shape=(), dtype=tf.string),  
})
```

- Examples (tfds.as_dataframe):

{% framebox %}

Display examples...

```
<script> const url = "https://storage.googleapis.com/tfds-data/visualization/dataframe/imdb_reviews-plain_text-1.0.0.html"; const dataButton  
= document.getElementById('displaydataframe'); dataButton.addEventListener('click', async () => { // Disable the button after clicking  
(dataframe loaded only once). dataButton.disabled = true;  
const contentPane = document.getElementById('dataframecontent'); try { const response = await fetch(url); // Error response codes don't throw  
an error, so force an error to show // the error message. if (!response.ok) throw Error(response.statusText);
```

```
const data = await response.text();  
contentPane.innerHTML = data;
```

```
} catch (e) { contentPane.innerHTML = 'Error loading examples. If the error persist, please open ' + 'a new issue.'; } }); </script>
```

{% endframebox %}

275 lines (209 sloc) | 9.65 KB

<> [icon] Raw Blame [icon] [icon] [icon]

imdb_reviews/bytes

- **Config description:** Uses byte-level text encoding with `tfds.deprecated.text.ByteTextEncoder`
- **Features:**

```
FeaturesDict({  
  'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=2),  
  'text': Text(shape=(None,), dtype=tf.int64, encoder=<ByteTextEncoder vocab_size=257>),  
})
```

- **Examples** (`tfds.as_dataframe`):

{% framebox %}

Display examples...

```
<script> const url = "https://storage.googleapis.com/tfds-data/visualization/dataframe/imdb_reviews-bytes-1.0.0.html"; const dataButton =  
document.getElementById('displaydataframe'); dataButton.addEventListener('click', async () => { // Disable the button after clicking (dataframe  
loaded only once). dataButton.disabled = true;  
const contentPane = document.getElementById('dataframecontent'); try { const response = await fetch(url); // Error response codes don't throw  
an error, so force an error to show // the error message. if (!response.ok) throw Error(response.statusText);
```

```
const data = await response.text();  
contentPane.innerHTML = data;
```

```
} catch (e) { contentPane.innerHTML = 'Error loading examples. If the error persist, please open ' + 'a new issue.'; } }); </script>
```

{% endframebox %}

imdb_reviews (enhanced)

275 lines (209 sloc) | 9.65 KB

<> [icon] Raw Blame [icon] [icon] [icon]

imdb_reviews/subwords8k

- Config description: Uses `tfds.deprecated.text.SubwordTextEncoder` with 8k vocab size
- Features:

```
FeaturesDict({  
  'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=2),  
  'text': Text(shape=(None,), dtype=tf.int64, encoder=<SubwordTextEncoder vocab_size=8185>),  
})
```

- Examples (`tfds.as_dataframe`):

{% framebox %}

Display examples...

```
<script> const url = "https://storage.googleapis.com/tfds-data/visualization/dataframe/imdb_reviews-subwords8k-1.0.0.html"; const  
dataButton = document.getElementById('displaydataframe'); dataButton.addEventListener('click', async () => { // Disable the button after  
clicking (dataframe loaded only once). dataButton.disabled = true;  
const contentPane = document.getElementById('dataframecontent'); try { const response = await fetch(url); // Error response codes don't throw  
an error, so force an error to show // the error message. if (!response.ok) throw Error(response.statusText);
```

```
const data = await response.text();  
contentPane.innerHTML = data;
```

```
} catch (e) { contentPane.innerHTML = 'Error loading examples. If the error persist, please open ' + 'a new issue.'; }); </script>
```

{% endframebox %}

```
import tensorflow_datasets as tfds
imdb, info = tfds.load('imdb_reviews/subwords8k', with_info=True, as_supervised=True)
```



```
train_data, test_data = imdb['train'], imdb['test']
```

```
tokenizer = info.features['text'].encoder
```

tensorflow.org/datasets/api_docs/python/tfds/features/text/SubwordTextEncoder

```
import keras_nlp

imdb = tfds.load("imdb_reviews", as_supervised=True, data_dir='./data', download=False)

train_reviews = imdb['train'].map(lambda review, label: review)
train_labels = imdb['train'].map(lambda review, label: label)

keras_nlp.tokenizers.compute_word_piece_vocabulary(
    train_reviews,
    vocabulary_size=8000,
    reserved_tokens=["[PAD]", "[UNK]"],
    vocabulary_output_file='imdb_vocab_subwords.txt'
)

subword_tokenizer = keras_nlp.tokenizers.WordPieceTokenizer(
    vocabulary='./imdb_vocab_subwords.txt'
)
```



```
import keras_nlp
```

```
imdb = tfds.load("imdb_reviews", as_supervised=True, data_dir='./data', download=False)
```

```
train_reviews = imdb['train'].map(lambda review, label: review)
```

```
train_labels = imdb['train'].map(lambda review, label: label)
```

```
keras_nlp.tokenizers.compute_word_piece_vocabulary(  
    train_reviews,  
    vocabulary_size=8000,  
    reserved_tokens=["[PAD]", "[UNK]"],  
    vocabulary_output_file='imdb_vocab_subwords.txt'  
)
```

```
subword_tokenizer = keras_nlp.tokenizers.WordPieceTokenizer(  
    vocabulary='./imdb_vocab_subwords.txt'  
)
```

```
import keras_nlp

imdb = tfds.load("imdb_reviews", as_supervised=True, data_dir='./data', download=False)

train_reviews = imdb['train'].map(lambda review, label: review)
train_labels = imdb['train'].map(lambda review, label: label)

keras_nlp.tokenizers.compute_word_piece_vocabulary(
    train_reviews,
    vocabulary_size=8000,
    reserved_tokens=["[PAD]", "[UNK]"],
    vocabulary_output_file='imdb_vocab_subwords.txt'
)

subword_tokenizer = keras_nlp.tokenizers.WordPieceTokenizer(
    vocabulary='./imdb_vocab_subwords.txt'
)
```

```
import keras_nlp

imdb = tfds.load("imdb_reviews", as_supervised=True, data_dir='./data', download=False)

train_reviews = imdb['train'].map(lambda review, label: review)
train_labels = imdb['train'].map(lambda review, label: label)

keras_nlp.tokenizers.compute_word_piece_vocabulary(
    train_reviews,
    vocabulary_size=8000,
    reserved_tokens=["[PAD]", "[UNK]"],
    vocabulary_output_file='imdb_vocab_subwords.txt'
)

subword_tokenizer = keras_nlp.tokenizers.WordPieceTokenizer(
    vocabulary='./imdb_vocab_subwords.txt'
)
```

```
import keras_nlp

imdb = tfds.load("imdb_reviews", as_supervised=True, data_dir='./data', download=False)

train_reviews = imdb['train'].map(lambda review, label: review)
train_labels = imdb['train'].map(lambda review, label: label)

keras_nlp.tokenizers.compute_word_piece_vocabulary(
    train_reviews,
    vocabulary_size=8000,
    reserved_tokens=["[PAD]", "[UNK]"],
    vocabulary_output_file='imdb_vocab_subwords.txt'
)

subword_tokenizer = keras_nlp.tokenizers.WordPieceTokenizer(
    vocabulary='./imdb_vocab_subwords.txt'
)
```



```
sample_string = 'TensorFlow, from basics to mastery'

tokenized_string = subword_tokenizer.tokenize(sample_string)
print('Tokenized string is {}'.format(tokenized_string))

original_string = subword_tokenizer.detokenize(tokenized_string).numpy().decode("utf-8")
print('The original string: {}'.format(original_string))
```



```
sample_string = 'TensorFlow, from basics to mastery'
```

```
tokenized_string = subword_tokenizer.tokenize(sample_string)  
print('Tokenized string is {}'.format(tokenized_string))
```

```
original_string = subword_tokenizer.detokenize(tokenized_string).numpy().decode("utf-8")  
print('The original string: {}'.format(original_string))
```



```
sample_string = 'TensorFlow, from basics to mastery'
```

```
tokenized_string = subword_tokenizer.tokenize(sample_string)
```

```
print('Tokenized string is {}'.format(tokenized_string))
```

```
original_string = subword_tokenizer.detokenize(tokenized_string).numpy().decode("utf-8")
```

```
print('The original string: {}'.format(original_string))
```



```
sample_string = 'TensorFlow, from basics to mastery'

tokenized_string = subword_tokenizer.tokenize(sample_string)
print('Tokenized string is {}'.format(tokenized_string))

original_string = subword_tokenizer.detokenize(tokenized_string).numpy().decode("utf-8")
print('The original string: {}'.format(original_string))
```

```
Tokenized string is [ 53 2235 543 1827 3024 13 198 1659 174 167 2220 238]
The original string: TensorFlow , from basics to mastery
```

```
for i in range(len(tokenized_string)):
    subword = subword_tokenizer.detokenize(tokenized_string[i:i+1]).numpy().decode("utf-8")
    print(subword)
```



```
for i in range(len(tokenized_string)):
    subword = subword_tokenizer.detokenize(tokenized_string[i:i+1]).numpy().decode("utf-8")
    print(subword)
```

```
T
##ens
##or
##F
##low
,
from
basic
##s
to
master
##y
```

```
embedding_dim = 64
model = tf.keras.Sequential([
    tf.keras.Input(shape=(MAX_LENGTH,)),
    tf.keras.layers.Embedding(subword_tokenizer.vocabulary_size(), EMBEDDING_DIM),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(6, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model.summary()
```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 120, 64)	488,460
global_average_pooling1d_1 ((None, 64)	0
dense_4 (Dense)	(None, 6)	390
dense_5 (Dense)	(None, 1)	7
Total params: 489,037		
Trainable params: 489,037		
Non-trainable params: 0		

```
num_epochs = 10
```

```
model.compile(loss='binary_crossentropy',  
              optimizer='adam',  
              metrics=['accuracy'])
```

```
history = model.fit(train_dataset,  
                    epochs=num_epochs,  
                    validation_data=test_data)
```